



SAN FRANCISCO STATE UNIVERSITY

College of Business · ISYS 812

Fintech Marketing Intelligence

Complaint Signal Analysis & Prescriptive Strategy

Group 2 · Hunsa Tiwana, Jason Tate, Marko Asaulyuk

Instructor: Dr. Lutfus Sayeed

May 2026

Data source

Consumer Financial Protection Bureau (CFPB) public complaint database

755,671 labeled rows · Temporal split: train 2014–2015 · test 2016–2017

EXECUTIVE SUMMARY

This project uses a multi-source customer behavior dataset with 777,867 total records. For the core analysis, we restrict to the CFPB subset (755,671 rows) to keep schema and outcome definitions consistent. The analysis asks where financial institutions fail, which failure patterns predict escalation, and which product opportunities those signals reveal.

We cleaned the dataset, answered five analytical questions, built an escalation prediction model, and turned those findings into product recommendations and a market gap analysis. The later sections go beyond the data. They apply findings to real product decisions, startup vs. enterprise strategy, and three market opportunities where no direct competitor was identified in our targeted scan.

The core finding:

Billing problems and unclear dispute processes are the two strongest escalation predictors in this dataset. The model built to detect them scores perfectly on a temporal holdout, though that result is dataset-specific, not a claim about how it would perform on different data. That same model, turned into a product, is the foundation of SignalStack, the highest-ranked opportunity in Part 7. No direct competitor was identified in our targeted scan.

How this report is organized

Section	Rubric requirement	Points
1. The Dataset	Name, column descriptions, shape, missing values	5 pts
2. Driving Question	Central question motivating the analysis	10 pts
3. Data Cleaning	Step-by-step cleaning rationale and first rows of clean data	10 pts
4. Data Analysis	Five questions with methods, findings, and visual references	10 pts
5. Data Modeling	Escalation prediction model, validation, feature analysis	10 pts
6-8. Prescriptive and Strategic	Beyond-rubric: product recommendations, market gaps, roadmap	bonus
9. Conclusion	Benefits gained and challenges encountered	5 pts
10. Limitations	What this analysis can and cannot claim	noted

PART 1 — DATASET INTRODUCTION

The Dataset

Dataset name

master_customer_behavior_clean.csv: a multi-source customer behavior dataset built for marketing intelligence analysis, with the CFPB (Consumer Financial Protection Bureau) as the primary source for this project.

What is in the dataset

Each row represents one customer event: a complaint filed with the CFPB, an app store review, a Google Trends data point, or a signal from another channel. The CFPB rows are the core of this analysis. The columns that matter most:

Column	What it contains
record_id	Unique identifier for each row
source	Where the record came from: CFPB, app store review, Google Trends, etc.
date / year / month	When the event was recorded
company	The financial institution or company the complaint or review is about
product_service	The financial product category: mortgage, credit card, debt collection, student loan, etc.
event_type	The type of event: customer complaint, review, or trend signal
text	The complaint or review text, or a keyword phrase for trend rows
sentiment_score / sentiment_label	Automated sentiment outputs from the data supplier. Not used as headline metrics for CFPB rows due to unreliable scoring on short boilerplate complaint text.
customer_action	The behavioral outcome label. Values include formal_complaint, complaining, churning, and others. This is the key outcome variable for the escalation model.
issue	The specific complaint issue parsed from CFPB metadata: billing disputes, debt collection conduct, account management failures, etc. Critical for both descriptive analysis and model features.
sentiment_revised_label	Project-added field that forces formal_complaint rows to negative, correcting misleading supplier labels on short complaint text.
url	Source link when available. Missing for some rows.
country	Geography when available. Missing for some rows.
raw_metadata	JSON metadata from the source, including complaint ID, company response, and submission channel.

Dataset shape and size

Metric	Value
Total rows	777,867
Total columns	17
CFPB rows (core analysis subset)	755,671
Time window (CFPB concentration)	2014 to 2017
Fintech proxy cohort: PayPal/Venmo	3,655 rows
Fintech proxy cohort: Square/Block	687 rows
Large-bank baseline	167,474 rows

Missing values by column

Twelve of the 17 columns have zero missing values. The three columns with gaps are:

Column	Missing values	Explanation
url	2,702	Source links are not available for all records and are not used in any analysis.
issue	22,196	Non-CFPB rows have no issue field. All CFPB rows used in this analysis are complete on this column.
country	26,937	Geography is not available across all data sources and is not used as a model feature.

Why missing values do not affect the core analysis:

The three columns with missing values are either non-essential (url is a source link only) or missing exclusively for non-CFPB rows. Since the escalation model and complaint analysis use only CFPB rows, all required fields are complete within that working subset.

PART 2 — DRIVING QUESTION

The Driving Question

Driving question:

How do CFPB complaint patterns (issue type, product line, company, and customer action) signal escalation risk, and which startup and enterprise opportunities does that signal point to?

Why this question

Financial companies generate enormous amounts of complaint data. Most treat it as a compliance cost: something to manage, not study. The signal never reaches the product team. Our question is whether that data is actually predictive and, if so, what to do with it.

It works at three levels. Descriptively: where do complaints concentrate? Predictively: can complaint text and product type predict whether a customer files formally? Prescriptively: if the signal holds up, what does a PM do with it?

How does the analysis answer it

Layer	What we did	Where it appears in this report
Descriptive	Mapped complaint volume, escalation rates, and issue patterns by product and company tier	Section 4, Questions 1-3
Predictive	Built a logistic regression model using complaint text and product type to predict formal escalation	Section 4, Question 4 and Section 5
Prescriptive (operational)	Translated the top signal clusters into four specific product feature recommendations	Section 6
Prescriptive (strategic)	Identified three market gaps (products that do not yet exist) that the signal points toward	Part 7

PART 3 — DATA CLEANING

Data Cleaning Process

All cleaning logic lives in `shared_cleaning_rules.py`: the one file every notebook, workflow, and training script reads from. Running it on the raw data always produces the same clean CSV, which is what keeps every result reproducible.

Step-by-step cleaning actions

Step	What was done	Why
1	Dropped the rating column	Not used in any analysis and absent from most source schemas. Removing it standardizes the table structure across all downstream operations.
2	Defined which columns are allowed to have empty values	URL, country, and <code>raw_metadata</code> are often blank for certain data sources. Keeping rows where these fields are empty avoids over-pruning valid complaint records from sources that simply do not provide this information.
3	Removed rows with missing values in required columns	Any row missing a value in core analytical columns (<code>company</code> , <code>product_service</code> , <code>customer_action</code> , <code>text</code> , etc.) was dropped. This ensures every analysis and model training step has complete inputs.
4	Parsed the issue field from <code>raw_metadata</code> JSON	CFPB-specific issue categories (billing disputes, debt collection attempts, account closures, etc.) are stored in a JSON metadata field, not a standalone column. This step extracts them into their own column, making the issue taxonomy available as both an analysis dimension and a model feature.
5	Created <code>sentiment_revised_label</code>	The supplier's sentiment labels are unreliable for short CFPB complaint text. Many formal complaints were labeled positive or neutral because the text is boilerplate. This step forces all <code>formal_complaint</code> rows to negative, producing consistent labels for notebook comparisons, while leaving the original sentiment fields untouched.
6	Created <code>trends_keyword</code> for Google Trends rows	Google Trends records store their search keyword in the <code>company</code> column rather than a dedicated field. This step parses the keyword out so the <code>company</code> column is not contaminated by trend phrases during company-level analysis.
7	Exported the standardized clean file	The cleaned dataset is written to <code>master_customer_behavior_clean.csv</code> . Every notebook, script, dashboard, and workflow reads from that one file. Everyone works with the same data.

First rows of the cleaned dataset

The table below shows the first five rows of master_customer_behavior_clean.csv across the six most analytically relevant columns. The full 17-column view is available in the Python notebook submission (descriptive/descriptive_analysis.ipynb).

record_id	source	company	product_service	customer_action	issue
1	CFPB	Rocket Mortgage, LLC	Mortgage	formal_complaint	Application, originator, mortgage broker
2	CFPB	SYNCHRONY FINANCIAL	Credit card	formal_complaint	Credit line increase/decrease
3	CFPB	GREAT LAKES	Student loan	formal_complaint	Can't repay my loan
4	CFPB	DISCOVER BANK	Student loan	formal_complaint	Dealing with my lender or servicer
5	CFPB	HYUNDAI CAPITAL AMERICA	Debt collection	formal_complaint	Cont'd attempts collect debt not owed

What the first rows tell us:

Every CFPB row is a complaint record linked to a specific company, product line, and parsed issue. Most rows in this extract have customer_action set to formal_complaint, the dominant label at roughly 97.8% of CFPB rows. A smaller set of other customer_action labels (closed without relief, in-progress resolutions) is used as the non-formal class in the escalation model. The model uses issue text and product category to distinguish between them.

PART 4 — DATA ANALYSIS

Data Analysis: Five Questions

Five questions, each answered using a different method, are ordered to build on one another.

Question 1: Which combinations of products and issues drive the highest complaint volume?

Method: We counted complaints by product type and issue category, then ranked them by volume. All counts use the CFPB-only subset of the dataset.

Key finding: Mortgage, debt collection, bank account/service, and credit reporting have the highest complaint volumes. Within those, the dominant issues are billing disputes, collection attempts on debts not owed, and servicing failures. These four products account for the majority of all formal CFPB complaints. Spread an intervention evenly across all product lines, and it loses its impact on the cases that actually need it.




Figure 1 — Complaint volume by product line (CFPB, top 12 categories)

Question 2: How does the distribution of formal complaints vary across customer action types?


Method: We counted how often each customer action type (formal complaint, in-progress, closed, etc.) appears in the CFPB data, then calculated the percentage that are formal complaints. We broke this down by product line to see where formal escalation concentrates.

Key finding: Formal complaints are approximately 97.8% of all CFPB rows. The remaining 2.2% are in-progress cases, closed-without-relief resolutions, and similar outcomes. That means the model isn't distinguishing between complainers and non-complainers. It's distinguishing formal escalation from other CFPB action types. That's a finer and more useful distinction.

Question 3: How do major financial players differ from smaller ones in what they get complained about?

Method: We filtered out companies that appeared too few times to draw reliable conclusions, then calculated each company's share of total complaints and analyzed which issue types dominate their profiles. Major players have the highest complaint counts; smaller players have thinner, less consistent patterns.

Key finding: Major financial players have concentrated complaint profiles. A small number of recurring issues account for most of their complaints. Smaller players show noisier, more varied patterns. The fintech proxy cohort (PayPal/Venmo and Square/Block-side) shows distinctly different patterns from the large-bank baseline: PayPal/Venmo complaints center on transaction reliability and fraud, while Square concentrates in debt collection conduct. Both differ from the mortgage-and-servicing pattern that dominates large banks. That's why this report keeps fintech and bank recommendations in separate lanes.



Question 4: Can complaint text and product type predict escalation to a formal CFPB filing?

Method: We built a logistic regression classifier using two inputs: the text of each complaint (converted to numerical word-importance scores using TF-IDF, a standard text analysis technique) and the product line (converted to a set of yes/no flags the model can read). We trained the model on 2014–2015 data and tested it on 2016–2017 data it had never seen, which is the more realistic measure of how it would perform in practice.

Key finding: The model achieves a perfect discrimination score on the temporal holdout. That result is strong but dataset-specific, not a universal performance guarantee. Escalation is predictable from complaint language and product type alone. Text carries most of the weight: removing the product type barely changes performance, while text alone scores just as well as the combined model. The top escalation-predicting terms are billing, fee, delinquent, disputes, and APR. Fraud and scam terms push the other way. Fraud cases need separate handling.

Question 5: Are the escalation signals stable enough to support product prioritization?

Method: Five tests were run to confirm that the model isn't just memorizing the training data. These covered: performance on future data it had never trained on, performance on product types it had never seen, how much accuracy drops when each input type is removed, whether the model's confidence scores match real outcomes, and whether adding generic complaint phrases changes the predictions.

Key finding: All five checks passed. The model generalizes to years and product types it was never trained on. Its confidence scores match the real outcome probabilities (calibration error: 0.0018, pass threshold: 0.08). Adding generic boilerplate to the complaint text doesn't shift its predictions. The signal is stable. Use it as a decision-support tool, not as a deterministic decision rule.

PART 5 — DATA MODELING

The Escalation Prediction Model

The model predicts whether a complaint escalates to a formal CFPB filing (Class 1) or takes some other path (Class 0). We trained on 2014–2015 data and tested on 2016–2017, a temporal holdout, which is the harder and more honest test of whether the model actually generalizes.

Model performance

Evaluation	ROC AUC	Avg Precision	n test rows	Notes
Random split (full model)	1.0	1.0	151,135	Upper bound reference
Temporal split (full model)	1.0	1.0	434,241	Primary generalization test
Product holdout (unseen products)	1.0	1.0	4,895	Holds on unseen product types
Product features only (ablation)	0.949	0.998	151,135	Text adds signal beyond product alone
Text features only (ablation)	1.0	1.0	151,135	Text is the primary driver

Classification report (temporal holdout)

Class	Precision	Recall	F1-Score	Support
Class 0: other action	1.000	0.747	0.855	9,703
Class 1: formal complaint	0.994	1.000	0.997	424,538
Weighted avg	0.994	0.994	0.994	434,241

Five independent checks — all passed

Check	Threshold	Result
Temporal generalization	ROC AUC $\geq$ 0.75	PASS — 1.0
Unseen product generalization	ROC AUC $\geq$ 0.70	PASS — 1.0
Incremental gain over product-only	Delta AUC $\geq$ 0.02	PASS — +0.051
Calibration usable	ECE $\leq$ 0.08	PASS — 0.0018
Not over-sensitive to canned phrases	p95 delta $\leq$ 0.20	PASS — 0.0005

What to make of ROC AUC 1.0:

A perfect score sounds suspicious, and it should. This reflects strong signal in this specific dataset. The complaints follow clear behavioral patterns. Don't read it as a universal law. All five validation checks passed, including temporal holdout and unseen product generalization, so the model isn't just memorizing training data. Use it as a decision-support tool, not ground truth.

The two-sided feature picture

Some complaint terms predict formal escalation. Others do the opposite: they predict cases that stay informal. The terms "fraud," "scam," and "identity theft" are the strongest predictors of non-formal outcomes, with "fraud" having nearly five times the predictive strength of any escalation-driver term. That's not noise. It confirms that fraud cases follow a behaviorally different path. Handling them through the same billing or dispute workflow as a fee complaint will lead to poor outcomes for both.

Direction	Top signal terms	Model weight range	Implication
→ Formal complaint	billing, delinquent, disputes, fee, APR, opening/closing	+0.51 to +1.98	These are the complaint escalation drivers
→ Other actions	fraud, scam, identity theft, cancelling account	-8.90 to -3.94	Fraud cases need a separate product path

PART 5 CONTINUED — FINTECH COHORT ANALYSIS

Where the Fintech Cohort Fails

The fintech cohort and the large-bank baseline fail in different ways. That matters. Recommendations built from mortgage and servicing data don't transfer cleanly to a digital payments company.

PayPal / Venmo: digital payment friction

3,655 rows. Most complaints cluster around money transfers and basic account management. The recurring issues: fraud and scams, transaction problems, funds not available when promised, and account access failures. Customers aren't complaining about bad UX. They're complaining that the money didn't move when it was supposed to, or that someone got into their account.

Square / Block-side: collections conduct

687 rows, heavily concentrated in debt collection. The pattern here is different: customers disputing debts they say they don't owe, collectors who can't verify the debt, communication tactics that cross legal lines. This isn't a payment reliability problem. It's a fairness and process problem.

Large-bank baseline: a different failure zone

167,474 rows dominated by mortgage servicing, foreclosure, and escrow issues. As a useful contrast, it shows why fintech-specific recommendations need their own lane, not a filtered version of the bank analysis.

The key finding here:

Fraud complaints in the PayPal/Venmo cohort don't behave like billing complaints. The model's negative-weight features make this explicit. You can't handle both with one workflow. A fraud case needs its own intake path, its own evidence handling, and its own resolution logic, not a shared queue with the fee dispute from the credit card team.

PART 6 — PRESCRIPTIVE RECOMMENDATIONS

What Fintech PMs Should Build Next

Four interventions, each tied to a specific signal cluster. None of them come from industry benchmarks. They come from the complaint data on what customers actually get stuck on.

1. Payment Assurance Layer

When a transfer is pending, customers don't know why or when it'll clear. Add a status trace, an expected availability timer, and exception flags when something goes wrong. That's the entire PayPal/Venmo funds-availability complaint cluster, addressed. KPI: payment-related complaint share down.

2. Fraud / Scam Triage Workflow

Fraud cases need their own front door. Separate intake, severity routing, an instant freeze option, SLA tracking. The model is unambiguous about this. Fraud cases are behaviorally distinct from billing disputes, and routing them into the same queue leads to poor outcomes for both. KPI: fraud case cycle time down, repeat fraud complaints down.

3. Collections Transparency Pack

For any firm with lending or collections exposure. Customers disputing debts they say they don't owe need a debt verification packet, clear communication controls, and an audit trail. This addresses the Square/Block complaint cluster directly. KPI: debt-verification complaint share down.

4. Account Lifecycle Guardrails

Opening, closing, and payoff moments generate a disproportionate share of complaints. The model flags these moments as significant predictors of escalation. Checklists, plain-language status transitions, and high-risk alerts at these moments reduce the confusion that drives escalation. KPI: account-management complaint recurrence down.

Track A vs Track B

Track	Timeline	Logic	Activation
Track A (primary)	0–3 months	Focus on the highest-priority interventions first. Ship quickly, measure, then expand.	Default. Start here.
Track B (fallback)	12+ months	Shift to a larger governance and monitoring build: cross-team coordination, risk dashboards, policy controls. Only activates if Track A isn't working.	2+ trigger conditions failing for 6–8 consecutive weeks

Track B trigger conditions

Switch to Track B review if two or more of these conditions fail for six to eight weeks running:

- Adoption stall: under 20% active use among the target test group at week 6
- No improvement in escalation rate: under 5% relative reduction vs the baseline
- Cycle-time miss: median resolution time not improving by at least 15%
- Recurrence non-improvement: repeat complaint rate flat or increasing
- Governance incident increase: compliance incidents up vs baseline month
- Economic deterioration: customer acquisition cost rising while the share of new users who stay and keep using the product declines

PART 7 — STRATEGIC GAP ANALYSIS

Strategic Gap Analysis: What the Market Is Missing

Two different questions:

The prescriptive section (Part 6) asks what existing players should do differently. Part 7 asks what product should exist that doesn't yet. The five concepts below aren't improvements to current offerings. They're responding to problems that existing financial companies can't fix without working against how they make money.

The Richard Branson approach

Richard Branson built Virgin Atlantic in 1984 after studying what major airlines were fundamentally bad at. Not temporarily. By design. They treated passengers as cargo. Branson didn't try to fix that from inside an existing airline. He built a new one whose entire culture inverted what the incumbents couldn't change without hurting their own business model. That same logic carried through to Virgin Mobile, Virgin Money, and the rest of the Virgin portfolio.

The same logic applies here. Financial institutions hold all the records, control what counts as a resolution, and show up to a CFPB dispute with a full database. The customer shows up with screenshots. Every complaint theme in this analysis traces back to that imbalance.

The Five Product Directions

#	Concept	Tagline	Gap it fills	Themes addressed
01	Financial Advocate	The first AI that negotiates on your behalf	Consumers have no proactive advocate when hit with fees or a denied dispute	1, 2, 7
02	Complaint Intelligence as a Service (ClaaS)	The complaint-risk intelligence layer for fintech operators	Complaint data sits in legal and ops. The PM and risk teams rarely see it	All 7
03	Trust Score Infrastructure	Score institutions the way institutions score consumers	When choosing a bank or fintech, consumers have no reliable data on how that institution actually handles complaints and disputes	4, 7 directly; 1, 2 indirectly
04	Embedded Compliance Co-Pilot	Compliance-by-design, embedded at build time	Every fintech team rebuilds the same compliance logic from scratch (fee disclosures, dispute handling, legal notices) and most do it inconsistently	1, 4, 7
05	Delinquency Prevention OS	Intervene before delinquency, not after	The pre-delinquency window (30–60 days before status change) has no product occupying it	3, 4, 5 directly; 2 indirectly

Three Unified Products

Three of the five concepts share core infrastructure, target the same user, or have complementary revenue models. Two natural clusters emerge. The fifth concept stands alone. For now.

Cluster A: TrustOS (Concepts 01 + 03)

The Trust Score tells a customer which institutions to avoid before they sign up. The Financial Advocate fights for them after they do. Together, they cover the full customer journey: pre-purchase intelligence and post-purchase muscle.

They run on the same infrastructure: CFPB complaint data, open banking connections (standardized APIs that let third-party apps securely access bank account data with customer permission), and a regulatory knowledge base covering UDAAP, TILA, and FCRA. The scoring data feeds the advocacy engine; the advocacy outcomes sharpen the scores.

Cluster B: SignalStack (Concepts 02 + 04)

The Compliance Co-Pilot catches compliance problems at build time, before bad product decisions become bad complaint patterns. ClaaS catches escalation risk at runtime, before internal cases become CFPB filings. Same customer (fintech PM, risk officer, compliance lead), different moments.

The model we built in this project powers both. They share a regulatory content library (UDAAP, TILA, and adverse action notices (the legally required letters sent when a customer is denied credit or service)) and the same audit log infrastructure (CFPB, DORA, and MiCA reporting formats). The go-to-market mirrors how Stripe grew. Stripe is the payment infrastructure company that scaled from developer tools into a full financial platform. A developer SDK (software toolkit embedded at build time) is adopted by engineering teams first, without requiring management sign-off, and the intelligence dashboard is the natural upgrade once it's in place.

Cluster C: PreventionOS (Concept 05, standalone)

PreventionOS doesn't fit neatly into either cluster. It requires deep account-level behavioral data that the other two don't. Its business model is B2B2C: the institution pays for it to reduce charge-offs, and the customer gets a financial health tool. That's different enough to build separately. Once it can show results, it's also the most obvious acquisition for either SignalStack or TrustOS.

Prioritization Framework

Five scoring dimensions, each pulled from this project's own analytical framework. Not generic startup criteria. These come from the same files that generated the operational recommendations.

Dimension	What it measures	Source
Signal coverage	Breadth of mapped complaint themes addressed	opportunity_themes.md + top features JSON
Time-to-market	Fit for rapid Track A execution	feature_and_gtm_strategy.md
Disruption delta	Degree of new-market creation vs incremental improvement	Strategic gap analysis criteria
Revenue defensibility	Switching costs, network effects, compounding moat	tradeoff_matrix.md
Regulatory tailwind	Whether current policy timelines support adoption now	Section 6 citations

Rank	Cluster	Signal	Speed	Disruption	Defensibility	Reg tailwind	Total
1	SignalStack (02+04)	5	4	4	5	5	23
2	TrustOS (01+03)	4	4	5	4	3	20
3	PreventionOS (05)	4	3	4	4	4	19

The reasoning behind each rank

SignalStack, Rank 1 (Score: 23). The escalation model we built in this analysis is the core intellectual property of this product. DORA (the EU's Digital Operational Resilience Act, effective January 2025) requires every financial firm to maintain tamper-proof audit logs. That regulation supports demand for exactly what SignalStack is built to provide. Once it's embedded in a client's development stack, removing it means rebuilding their compliance infrastructure from scratch. Unlike consumer products, B2B (business-to-business) revenue is predictable and contractual. We didn't just theorize about this product. We built its core engine as part of this analysis.

TrustOS, Rank 2 (Score: 20). The highest disruption potential of the three. It actually inverts how information flows in financial services. The honest caveat is that monetization is harder in the near term. Both consumers and institutions need to change behavior for it to work, and institutions have reasons to resist. The long-run moat is real, but it takes time to build. Edelman 2025 data is useful context here: financial services is trusted in only 17 of 28 countries at 64% global trust. There is a consumer demand for this.

PreventionOS, Rank 3 (Score: 19). Strong coverage on the delinquency and lifecycle themes. The gap it fills (the 30 to 60 day window before accounts go delinquent) is real and unoccupied. The lower rank reflects a practical constraint: it needs account-level behavioral data to prove its value, and getting that access takes time. We'd build it with the expectation that it would eventually be acquired by SignalStack or TrustOS, rather than growing into a fully independent platform.

Product Roadmap

SignalStack: roadmap

Track A (0–3 months)	Growth (3–12 months)	Track B (12+ months)
Compliance Co-Pilot SDK v1	ClaaS dashboard launch	Full compliance platform suite
Fee disclosure component	Escalation scoring by product × issue	Expanded regulator-report outputs
UDAAP content linter	Industry benchmarking	Advanced tamper-evident logging
Dispute workflow module	Alert threshold tuning	Cross-institution intelligence layer
Audit log API	CFPB report export	Real-time regulator dashboard

TrustOS: roadmap

Track A (0–3 months)	Growth (3–12 months)	Track B (12+ months)
Trust Score MVP	Advocate module launch	Full consumer trust platform
CFPB score aggregation	AI dispute drafting	Open-banking expansion
Consumer comparison layer	SLA timeline support	Tamper-evident score publication
Institution score v1	Institution benchmark dashboard	Policy / partner channel expansion

PreventionOS: roadmap

Pre-build (Track A parallel)	MVP (3–9 months)	Scale / Converge (12+ months)
Account data partnerships	30-day delinquency risk signal	Charge-off reduction validation
Behavioral model prototyping	Intervention options UI	Integration into hardship programs
Status messaging design	Plain-language alerts	Convergence with TrustOS or SignalStack

On blockchain:

Blockchain isn't the product. It's a feature that becomes worth using at two specific points: SignalStack at 12+ months, where DORA-era resilience and auditability requirements increase the value of tamper-evident logging features, supporting demand for on-chain record storage; and TrustOS at 12+ months, where institution scores recorded on a blockchain become legally defensible in a way that a regular database (which can be edited) is not. Three regulations cleared the path: GENIUS Act (July 2025, U.S. federal stablecoin framework), MiCA (December 2024, EU crypto-asset licensing), and DORA (January 2025, EU financial resilience mandate). Before 2025, using blockchain in a regulated financial context was a legal headache. Now there is a framework.

Competitor Landscape

We ran three targeted competitor scans, one per cluster. No direct competitor came up in this scan. The search was not exhaustive.

Cluster	Closest known players	Key differentiator	Gap status
TrustOS	Trustpilot, TrustScore.info, Newsweek/Statista rankings	Verified CFPB complaint-behavior scoring vs self-selected opinions; institution accountability vs platform marketing	None found in this scan
SignalStack	Sardine, Feedzai, Oscilar, Ballerine	Complaint-escalation intelligence vs fraud/crime detection; CFPB signal themes vs transaction anomaly signals	None found in this scan
PreventionOS	Credit Karma, budgeting apps, debt collectors	Pre-delinquency intervention (30–60-day window) vs post-hoc response; account-specific vs generic	None found in this scan

PART 8 — MARKET AND REGULATORY CONTEXT

Market Size and Why the Timing Is Right

Market size estimates

The TAM figures below are third-party research estimates. Treat them as directional rather than definitive. They vary by how the market is defined and shouldn't be cited as regulatory-grade data.

Market	2024–2025 estimate	Forecast	Source
B2B compliance SaaS / RegTech	\$15.24B (2025)	\$111.32B by 2034	Research and Markets, Global RegTech Market Report
Consumer fintech / Personal Finance Apps	\$132.92B (2024) → \$165.9B (2025)	\$406.5B by 2029	The Business Research Company

Regulatory citations

Regulation	Key date	Relevance to this analysis	Source
GENIUS Act	Enacted July 18, 2025	Federal stablecoin/on-chain record framework. Enables blockchain audit trail feature in SignalStack and TrustOS	U.S. Treasury: home.treasury.gov/news/press-releases/sb0197
MiCA	Full application from Dec 30, 2024	EU-passportable crypto-asset licensing. Enables on-chain features for EU-market roadmap	ESMA: esma.europa.eu
DORA (EU Digital Operational Resilience Act, effective Jan 2025)	Applies from Jan 17, 2025	Supports stronger auditability and operational resilience requirements across financial entities. SignalStack's compliance API is designed for this.	ESMA: esma.europa.eu
EU AI Act	High-risk obligations from Aug 2, 2026	Transparency and auditability required for AI-assisted financial decisions. Tailwind for Compliance Co-Pilot	EU AI Act Service Desk: ai-act-service-desk.ec.europa.eu
Edelman Trust Barometer 2025	Published 2025	Financial services trusted in only 17 of 28 countries; 64% global trust. Anchors TrustOS market narrative	edelmansmithfield.com/trust/2025

PART 9 — CONCLUSION

Conclusion

What we gained from this project

The most direct benefit was learning how to build and validate a supervised classification model on a large, imbalanced dataset, and making it hold up under scrutiny. The temporal split (training on 2014–2015 and testing on 2016–2017) forced us to think honestly about generalization, rather than chasing accuracy on data the model already knew. Working through the five validation checks also built a clearer intuition for the difference between a model that memorizes and one that actually learns a pattern.

The model was only part of what we learned. The harder skill was turning data results into actual decisions. Turning a logistic regression feature weight into a product recommendation requires judgment about causality, audience, and stakes. Doing that for a real domain (fintech complaints) was more useful than any generic machine learning exercise.

Going outside the dataset to study the market, regulations, and competitors added something we didn't expect to be this useful. The three product clusters (TrustOS, SignalStack, PreventionOS) came directly from asking what the complaint signal is demanding. Not from a separate market research pass. That's a method worth keeping: let the data suggest the market gaps, then go verify them.

Challenges we encountered and how we handled them

The first challenge was the data's age. The CFPB extract concentrates on 2014–2017, which is older than you'd want for a market-facing analysis. We chose internal consistency over recency. The dataset is analytically clean and comparable within its time window, and we clearly stated in the report that any new data validation should come before anyone acts on what we found.

The second challenge was the near-perfect model score. ROC AUC 1.0 on a holdout set sounds like overfitting, so we treated it as a warning sign and ran additional validation rather than calling it a win. The five validation checks (unseen product holdout, calibration test, phrase sensitivity test) were our direct response to that. Each check targeted a specific failure mode. All five passes were what gave us enough confidence to move forward with it.

The third challenge was Chime's absence from the dataset. Chime is one of the most visible fintech brands in consumer finance, but it doesn't appear by name in this CFPB extract. Rather than pretend the gap didn't exist or exclude fintech analysis entirely, we built a proxy cohort from fintech-adjacent firms that do appear and were explicit about what it represents. Being upfront about that gap from the start, building a clear workaround, and documenting it transparently is something we'd build into the project plan from day one next time.

The honest takeaway:

Good data analysis isn't about finding clean answers. It's about knowing exactly what your data can and cannot claim. Every limitation in this project (the data's age, the near-perfect model, the missing Chime records) pushed us to be more precise. That precision is the main thing this project built.

PART 10 — LIMITATIONS

What This Analysis Can and Cannot Claim

Every finding in this report comes with these caveats.

- This is a complaint dataset, not a customer survey. CFPB filers are people who got frustrated enough to file a formal complaint. They're not a random sample of all customers or all experiences.
- Chime doesn't appear in this extract. The proxy cohort (PayPal/Venmo, Square/Block, Upstart) fills that gap. Conclusions about fintech behavior reflect those companies specifically.
- ROC AUC 1.0 is dataset-specific. The model generalizes well within this extract: all five validation checks passed. It doesn't generalize universally, nor does it establish that billing friction causes escalation. Correlation, not causation.
- The TAM numbers are estimates. Market size figures vary significantly depending on how the analyst defines the market. Use them as order-of-magnitude context, not forecasts.
- The strategic concepts are hypotheses, not business plans. The five product directions and three clusters come from complaint signal analysis. They haven't been tested in the market, and the competitor scan isn't exhaustive.
- The data is from 2014–2017. The patterns are internally consistent and analytically sound. Before committing budget to any of the operational recommendations, validate against a more recent CFPB export.

What this analysis actually is:

It's a view of where specific fintech products failed customers in the CFPB complaint channel over a defined time window. Useful for product positioning, gap identification, and strategy. Not population sentiment, not demographic analysis, not revenue prediction.

SOURCES AND DATA FILES

Sources

Internal project files

- predictive/output_escalation_model/metrics.json
- predictive/output_escalation_model/classification_report.txt
- predictive/output_escalation_model/top_features_toward_formal_complaint.json
- predictive/output_escalation_model/top_features_toward_other_actions.json
- predictive/output_escalation_validation/validation_report.json
- prescriptive/opportunity_themes.md
- prescriptive/feature_and_gtm_strategy.md
- prescriptive/tradeoff_matrix.md
- prescriptive/fallback_trigger_rules.md
- prescriptive/fintech_competitive_slice.md
- prescriptive/prioritized_backlog.csv
- prescriptive/prescriptive_strategic_gap_analysis.md

Branding Standards

San Francisco State University. (n.d.). Brand guidelines. <https://brand.sfsu.edu/>

External citations — APA 7th edition

- Consumer Financial Protection Bureau. (2024). Consumer complaint database [Data set]. CFPB. <https://www.consumerfinance.gov/data-research/consumer-complaints/>
- U.S. Department of the Treasury. (2025, July 18). GENIUS Act — stablecoin legislation enactment communication. U.S. Treasury. <https://home.treasury.gov/news/press-releases/sb0197>
- European Securities and Markets Authority. (2024). Markets in Crypto-Assets Regulation (MiCA): Article 149 — entry into force and application. ESMA Interactive Single Rulebook. <https://www.esma.europa.eu/publications-and-data/interactive-single-rulebook/mica/article-149-entry-force-and-application>
- European Securities and Markets Authority. (2025). Digital Operational Resilience Act (DORA): Applicability and scope. ESMA. <https://www.esma.europa.eu/esmas-activities/digital-finance-and-innovation/digital-operational-resilience-act-dora>
- European Commission, AI Act Service Desk. (2026). EU AI Act implementation timeline. European Commission. <https://ai-act-service-desk.ec.europa.eu/en/ai-act/eu-ai-act-implementation-timeline>
- Edelman Smithfield. (2025). 2025 Edelman Trust Barometer: Financial services sector report. Edelman Smithfield. <https://www.edelmansmithfield.com/trust/2025/trust-barometer/report-financial-sector>
- Research and Markets. (2025). Global RegTech market report and forecast. Research and Markets. <https://www.researchandmarkets.com/reports/5926259/global-regtech-market-report-forecast>
- The Business Research Company. (2025). Personal finance apps global market report. The Business Research Company. <https://www.giiresearch.com/report/tbrc1705609-personal-finance-apps-global-market-report.html>

AI-assisted tools (Cursor, Claude) and Grammarly were used for coding support, brainstorming, and proofreading. All analyses and conclusions were developed and validated by the authors.